

Artificial Neural Networks

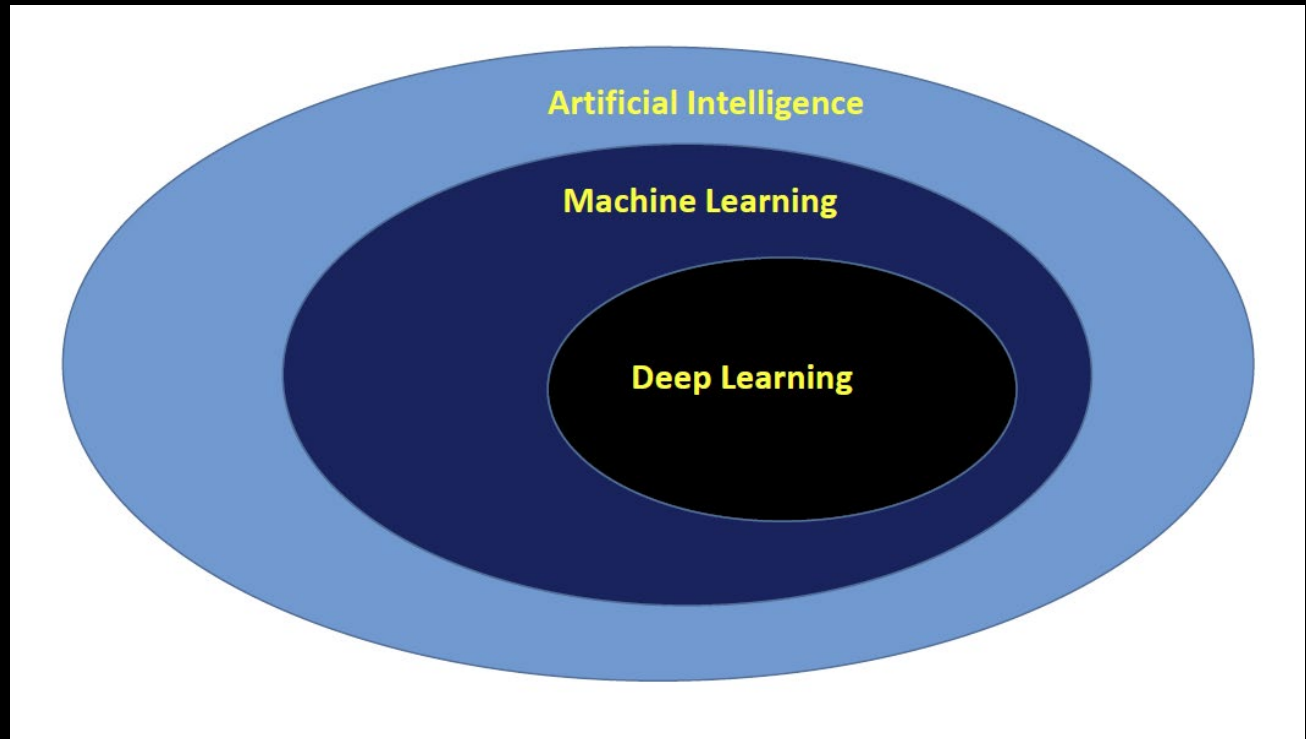
Chapter 10: pp 279 – 294, 323 – 329

Artificial Neural Networks (ANNs)


- Versatile, powerful, scalable
- The essence of Deep Learning
- Used to tackle highly complex ML tasks
 - Speech recognition
 - Object detection
 - Self-driving car
 - Recommendation systems
 - Machine translation
 - ...

- Popular Deep Learning Terms


- Deep Learning = Neural Networks with multiple hidden layers
- DNN = Deep Neural Network



From Biological to Artificial Neurons



- **Neuron:** computational building block for the brain
- Human brain:
 - ~100-1,000 trillion synapses

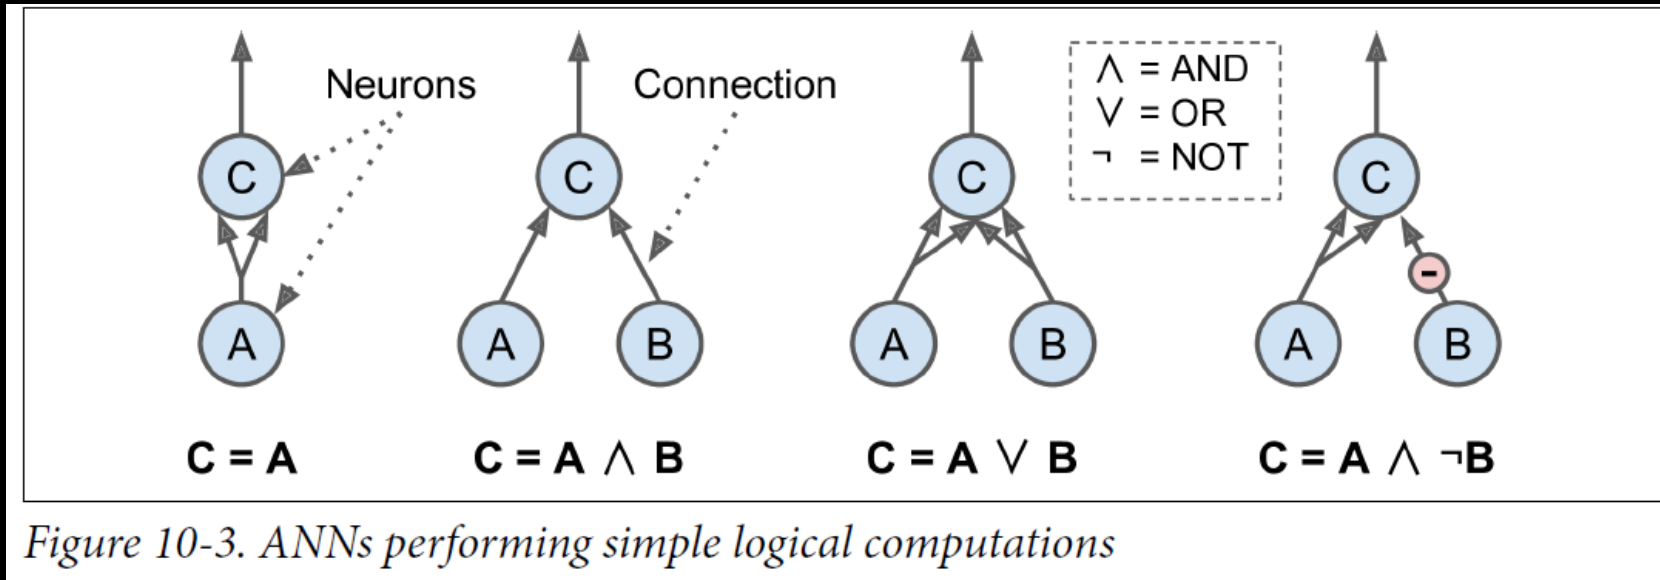


- **(Artificial) Neuron:** computational building block for the “neural network”
- (Artificial) neural network:
 - ~1-10 billion synapses

Human brains have ~10,000 computational power than computer brains.


Logic Computations with Neurons

- Simple logical computations



Perceptron (Rosenblatt's Perceptron)

- a.k.a Threshold Logic Unit (TLU), Linear Threshold Unit (LTU)
 - Inputs & outputs are real numbers instead of binary values
 - Each input connection is weighted
 - Output is the result of a step function of sum of the weighted inputs



Equation 10-1. Common step functions used in Perceptrons

$$\text{heaviside}(z) = \begin{cases} 0 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases} \quad \text{sgn}(z) = \begin{cases} -1 & \text{if } z < 0 \\ 0 & \text{if } z = 0 \\ +1 & \text{if } z > 0 \end{cases}$$

Figure 10-4. Threshold logic unit

- A single TLU can do simple linear binary classification
- A perceptron is a single layer *fully-connected* TLU's
- Every neuron is connected to the inputs

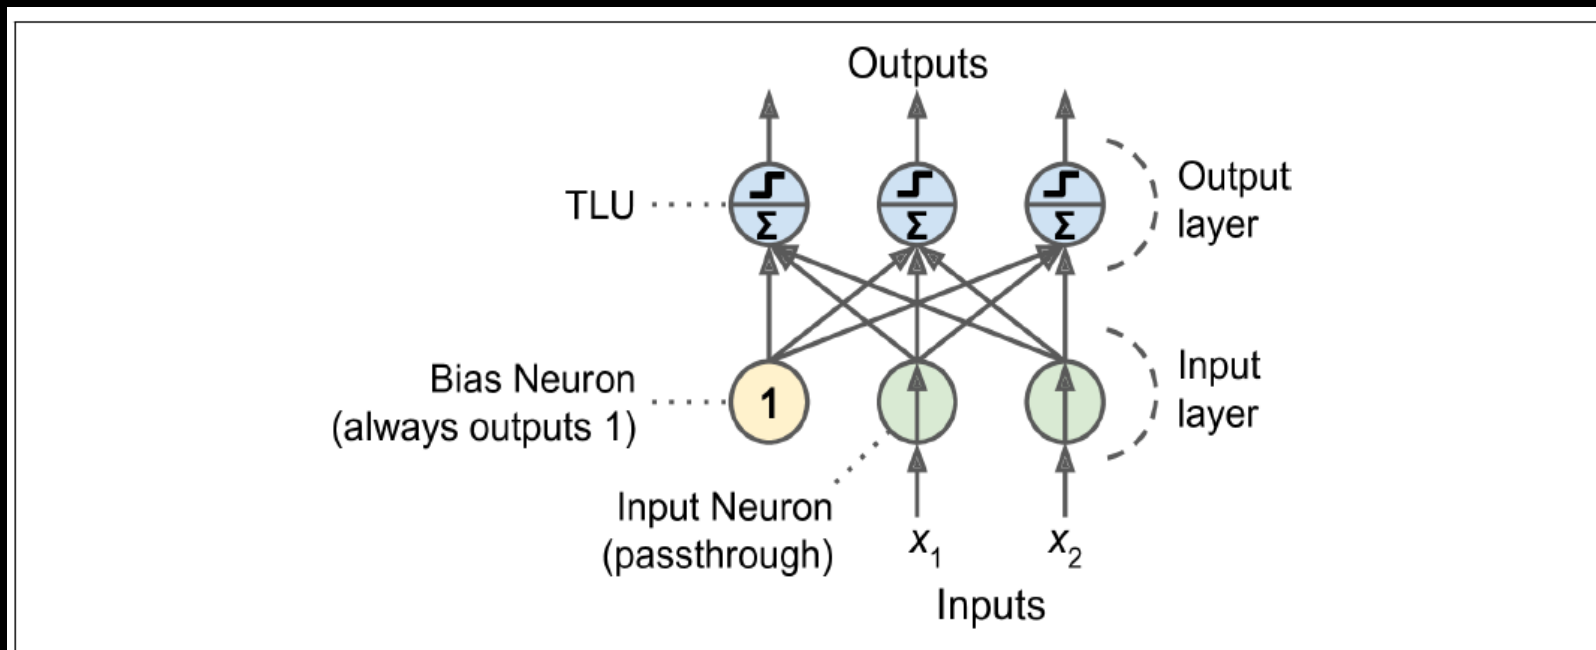


Figure 10-5. Perceptron diagram

- **X** is the input feature matrix.
 - One row per instance, one column per feature
- **W** is the connection weight matrix
 - Contains all connection weights except for the ones from bias neuron
 - One row per input, one column per neuron in the layer
- **b** is the bias vector
 - Contains all connection weights between the bias neuron and the other neurons
 - One bias term per neuron
- ϕ is the activation function
 - A step function

Training Perceptrons


- Hebb's Rule: the connection weight between two neurons is increased if they have the same output

Equation 10-3. Perceptron learning rule (weight update)

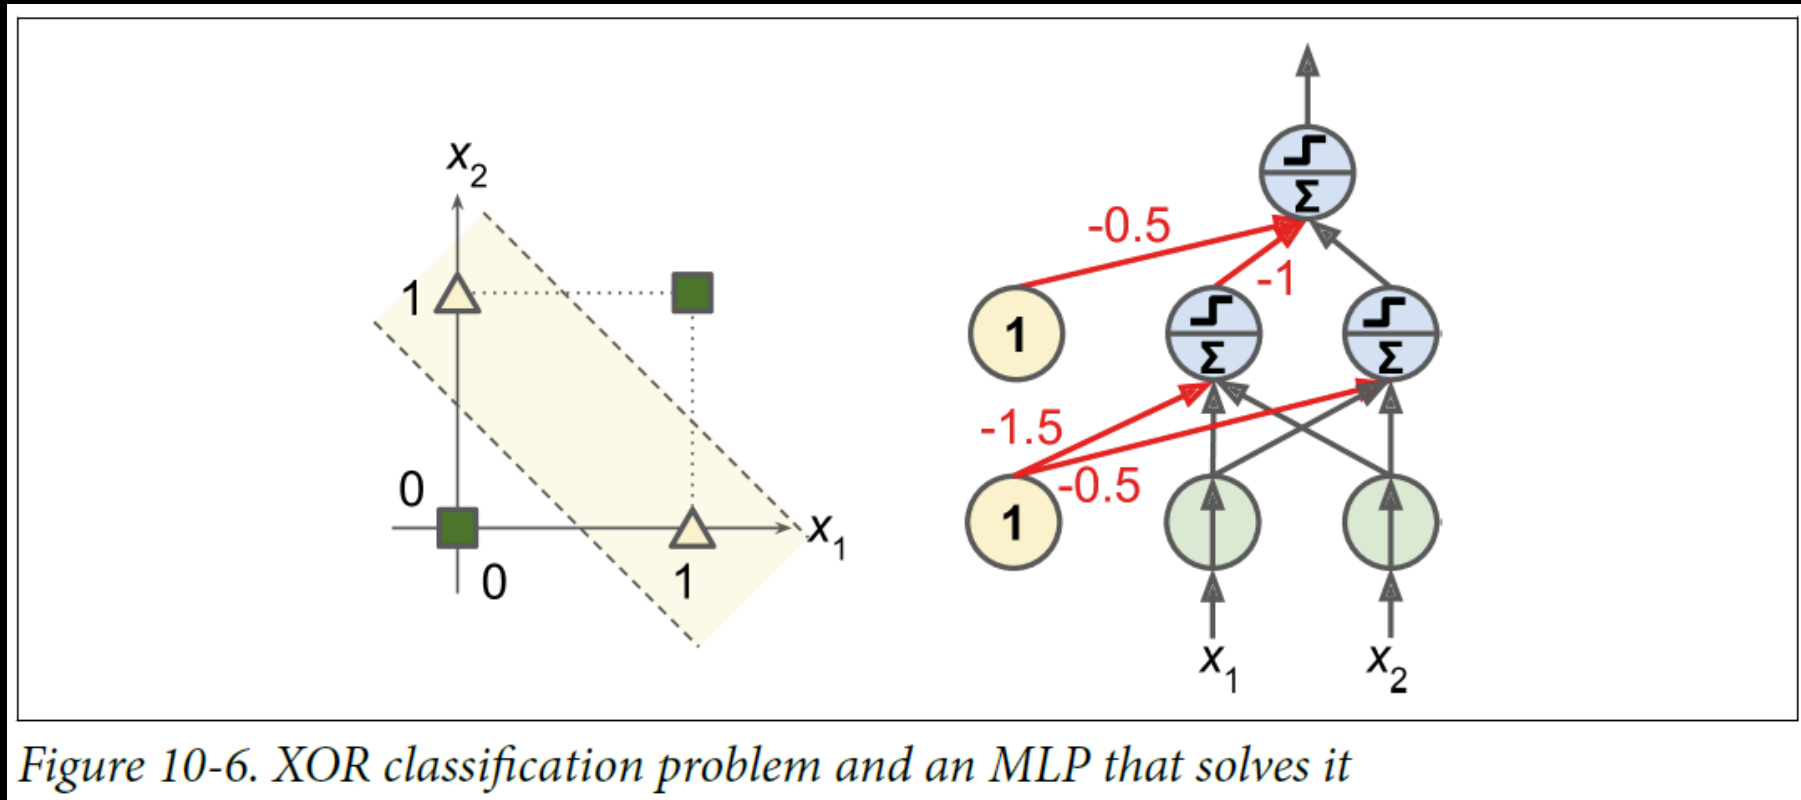
$$w_{i,j}^{(\text{next step})} = w_{i,j} + \eta(y_j - \hat{y}_j)x_i$$

- $w_{i,j}$ is the connection weight between the i^{th} input neuron and the j^{th} output neuron.
- x_i is the i^{th} input value of the current training instance.
- \hat{y}_j is the output of the j^{th} output neuron for the current training instance.
- y_j is the target output of the j^{th} output neuron for the current training instance.
- η is the learning rate.

- Perceptron Convergence Theorem
 - Algorithm will converge well for linearly separable classes
- Comparison with Logistic Regression
 - Perceptron outputs a class label based on a hard threshold
 - Logistic Regression outputs class probability

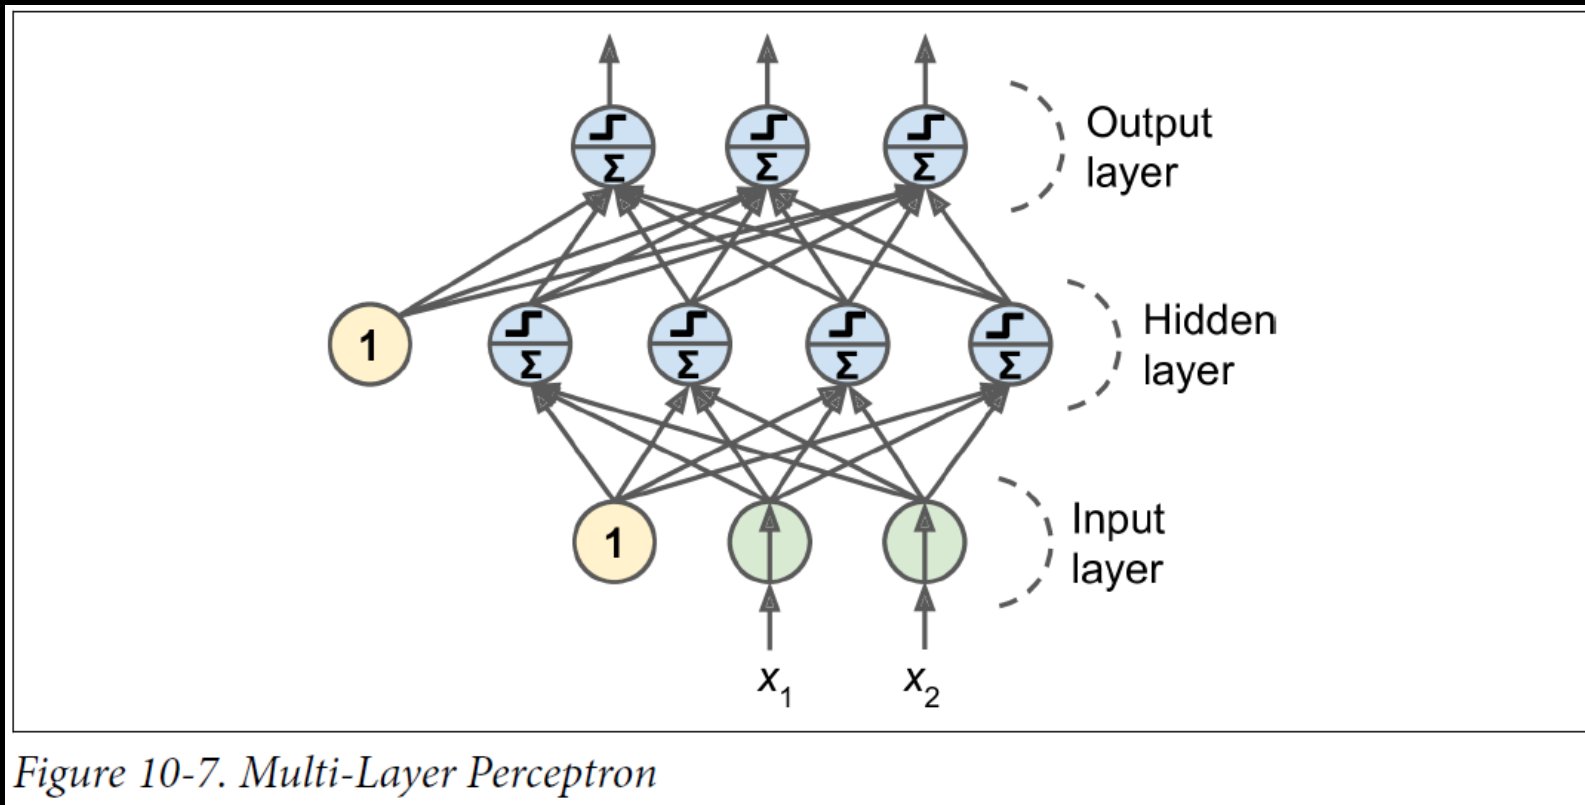


- Perceptron cannot solve XOR classification problem but Multi-Layer Perceptron (MLP) can



Multi-Layer Perceptron


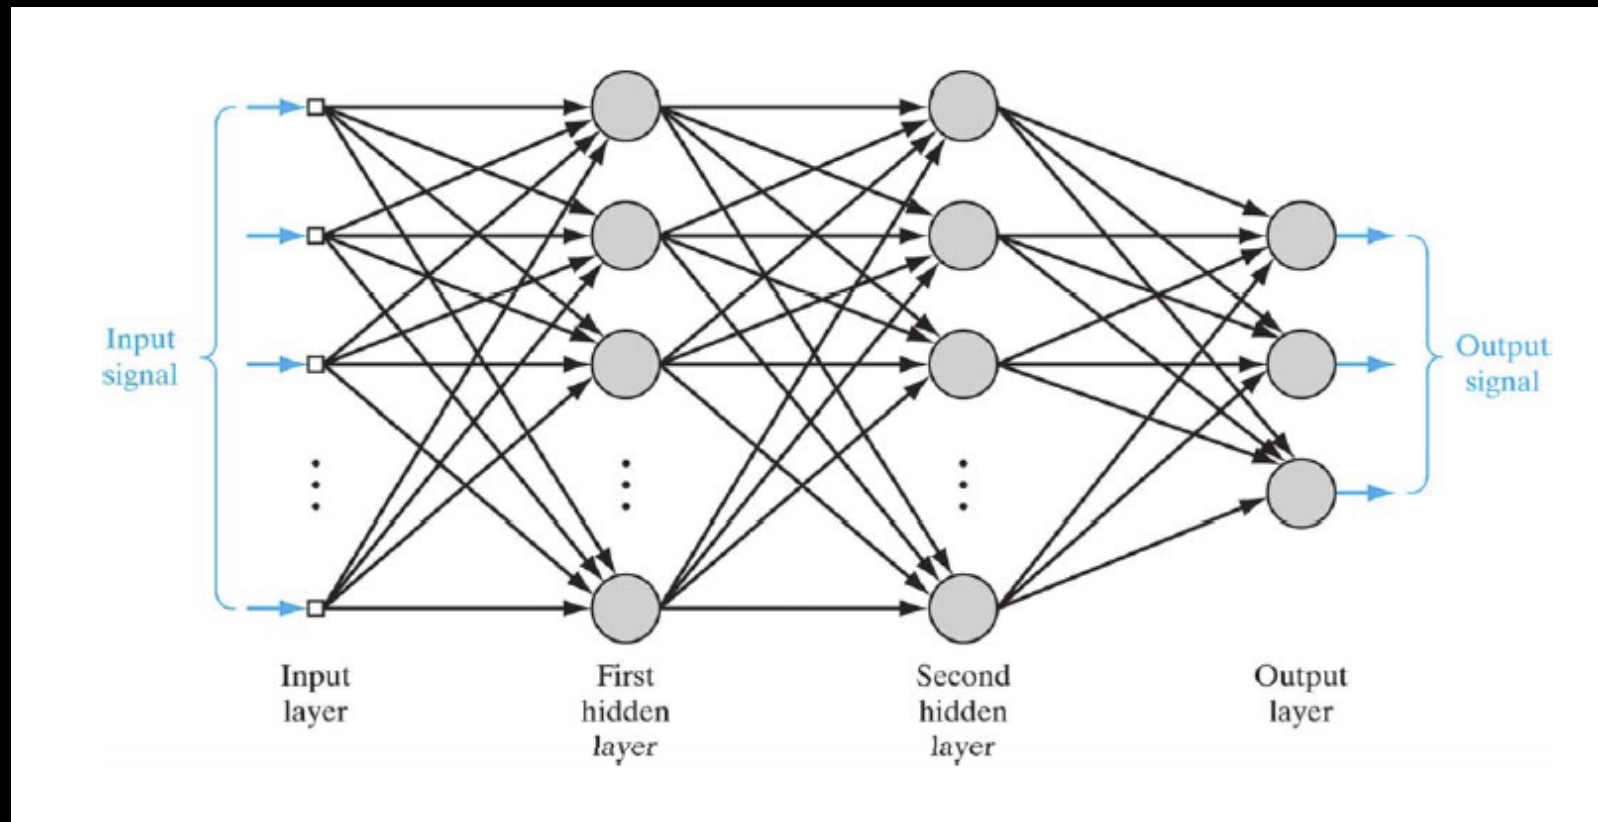
- MLP consists of input layer, at least one hidden layer, output layer
- Signal flows from inputs to outputs → feedforward NN
- Deep Neural Network (DNN) has many hidden layers



Function of the Hidden Neurons

- The hidden neurons act as *feature detectors*.
- As the learning process progresses across the multilayer perceptron, the hidden neurons begin to gradually “discover” the salient features that characterize the training data by performing a nonlinear transformation on the input data into a new space called *feature space*.
- In the feature space, the classes of interest in a pattern-classification task, for example, may be more easily separated from each other than could be in the original input data space.


Architectural graph of MLP with 2 hidden layers



- This is called a *three-layer* neural network
- # of input connections = # features
- # output neurons = # dimensions (regression) or # of labels/classes (classification)

Backpropagation Algorithm

- How do you train MLP?
 - Automatically find the error gradients with the training data and adjust the connection weights
- Forward pass
 - Makes predictions
 - Measures the error
- Reverse pass
 - Goes through each layer in reverse
 - Measures the error contribution from each connection
 - Modifies connection weights layer by layer to reduce the error



The Back-Propagation Algorithm

For an output neuron j ,

$$v_j(n) = \sum_{i=0}^m w_{ji}(n) y_i(n) \quad y_j(n) = \varphi(v_j(n))$$

$$\mathcal{E}(n) = \frac{1}{2} \sum_j e_j^2(n) = \frac{1}{2} \sum_j [d_j(n) - y_j(n)]^2$$

$$\begin{aligned} \frac{\partial \mathcal{E}(n)}{\partial w_{ji}(n)} &= \frac{\partial \mathcal{E}(n)}{\partial e_j(n)} \frac{\partial e_j(n)}{\partial y_j(n)} \frac{\partial y_j(n)}{\partial v_j(n)} \frac{\partial v_j(n)}{\partial w_{ji}(n)} \\ &= -e_j(n) \varphi'(v_j(n)) y_i(n) \end{aligned}$$

$$\Delta w_{ji}(n) = -\eta \frac{\partial \mathcal{E}(n)}{\partial w_{ji}(n)} = \eta \delta_j(n) y_i(n) \quad \text{where } \delta_j(n) = e_j(n) \varphi'(v_j(n))$$




FIGURE 4.3 Signal-flow graph highlighting the details of output neuron j .

- Neural Networks and Learning Machines by Simon Haykin

For a hidden neuron j ,

$$\mathcal{E}(n) = \frac{1}{2} \sum_k e_k^2(n) = \frac{1}{2} \sum_k \left[d_k(n) - \varphi \left(\sum_j w_{kj}^{out}(n) y_j^h(n) \right) \right]^2$$

$$y_j^h(n) = \varphi \left(\sum_i w_{ji}^h(n) y_i^{in}(n) \right)$$

$$\begin{aligned} \frac{\partial \mathcal{E}(n)}{\partial w_{ji}(n)} &= \sum_k e_k \frac{\partial e_k(n)}{\partial v_k(n)} \frac{\partial v_k(n)}{\partial y_j^h(n)} \frac{\partial y_j^h(n)}{\partial w_{ji}(n)} \\ &= - \sum_k e_k \varphi'(v_k(n)) w_{kj}(n) \varphi'(v_j(n)) y_i(n) \\ &= - \sum_k \delta_k(n) w_{kj}(n) \varphi'(v_j(n)) y_i(n) \end{aligned}$$

$$\Delta w_{ji}(n) = \eta \delta_j(n) y_i(n) \quad \text{where} \quad \delta_j(n) = \sum_k \delta_k(n) w_{kj}(n) \varphi'(v_j(n))$$

Backpropagated gradient
from the layer to the
immediate right

$$\begin{pmatrix} \text{weight} \\ \text{correction} \\ \Delta w_{ji}(n) \end{pmatrix} = \begin{pmatrix} \text{learning} \\ \text{rate} \\ \eta \end{pmatrix} \times \begin{pmatrix} \text{overall} \\ \text{local gradient} \\ \delta_j(n) \end{pmatrix} \times \begin{pmatrix} \text{input signal} \\ \text{of neuron } j \\ y_i(n) \end{pmatrix}$$




FIGURE 4.4 Signal-flow graph highlighting the details of output neuron k connected to hidden neuron j .


- Neural Networks and Learning Machines by Simon Haykin

(1) Introduction to Deep Learning with neon

Practical example: recognition of handwritten



MNIST dataset
70,000 images (28x28 pixels)
Goal: classify images into a digit 0-9



nervana


13:31 / 39:47



- Reference: Nervana/Intel

(1) Introduction to Deep Learning with neon

Practical example: recognition of handwritten



Total parameters:

$$W_{i \rightarrow j} \quad 784 \times 100$$

$$b_j \quad 100$$

$$W_{j \rightarrow k} \quad 100 \times 10$$


$$b_k \quad 10$$



nervana

(1) Introduction to Deep Learning with neon

Training procedure



1. Randomly seed weights
2. Forward-pass
3. Cost
4. Backward-pass




15:28 / 39:47



nervana

(1) Introduction to Deep Learning with neon

Forward pass

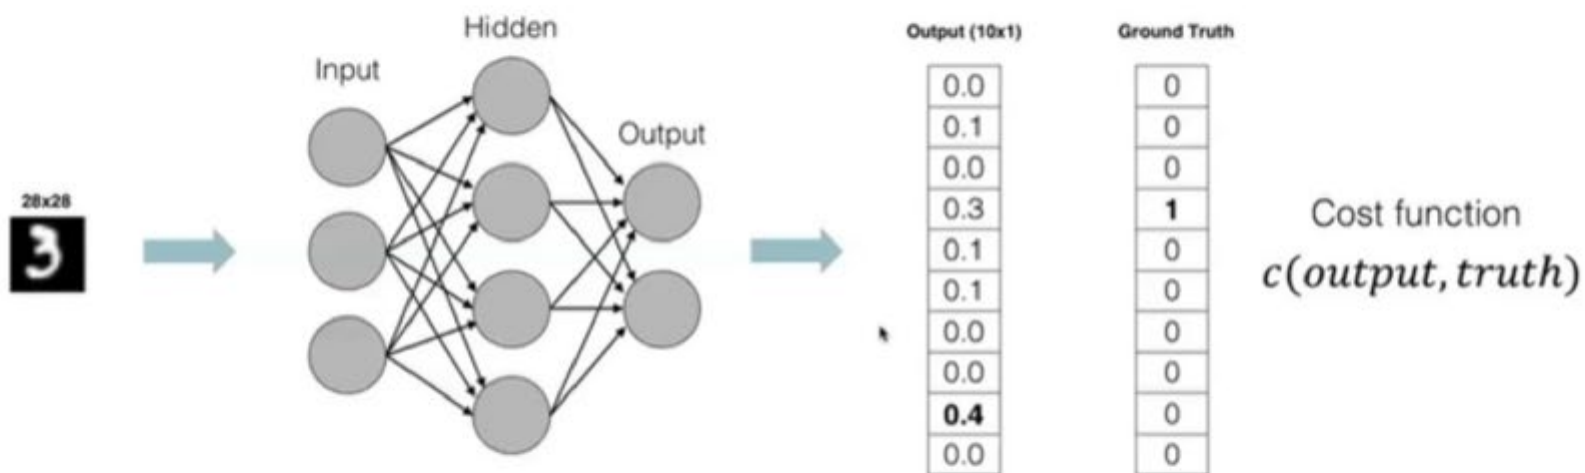


$$y = g \left(\sum_j W_j x_j + b \right)$$

nervana

(1) Introduction to Deep Learning with neon

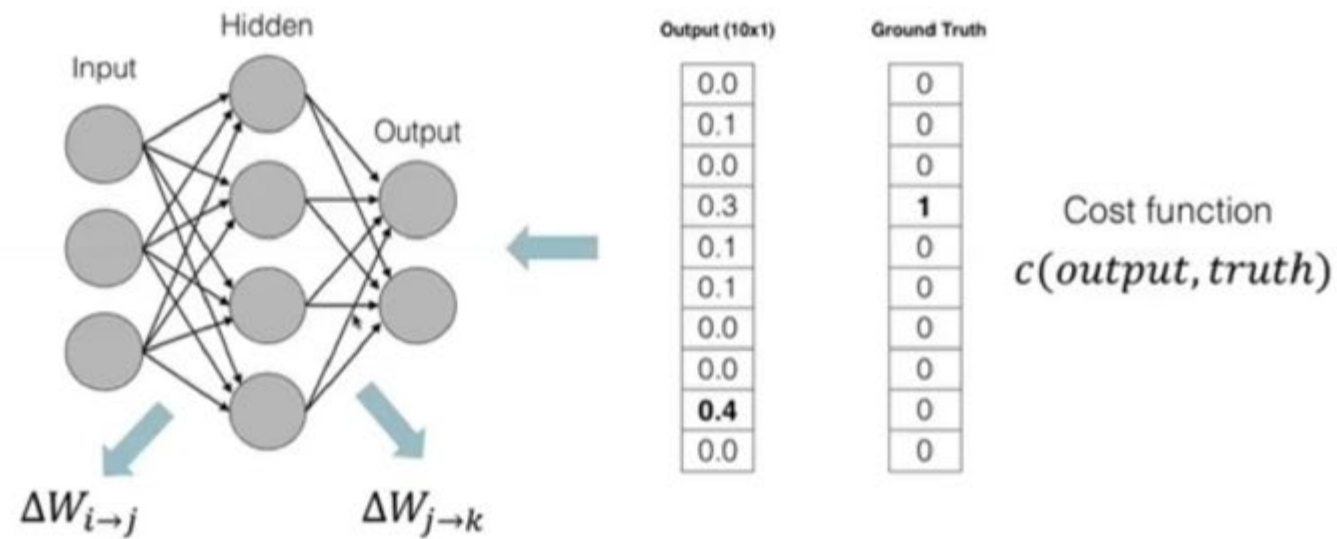
Cost



nervana

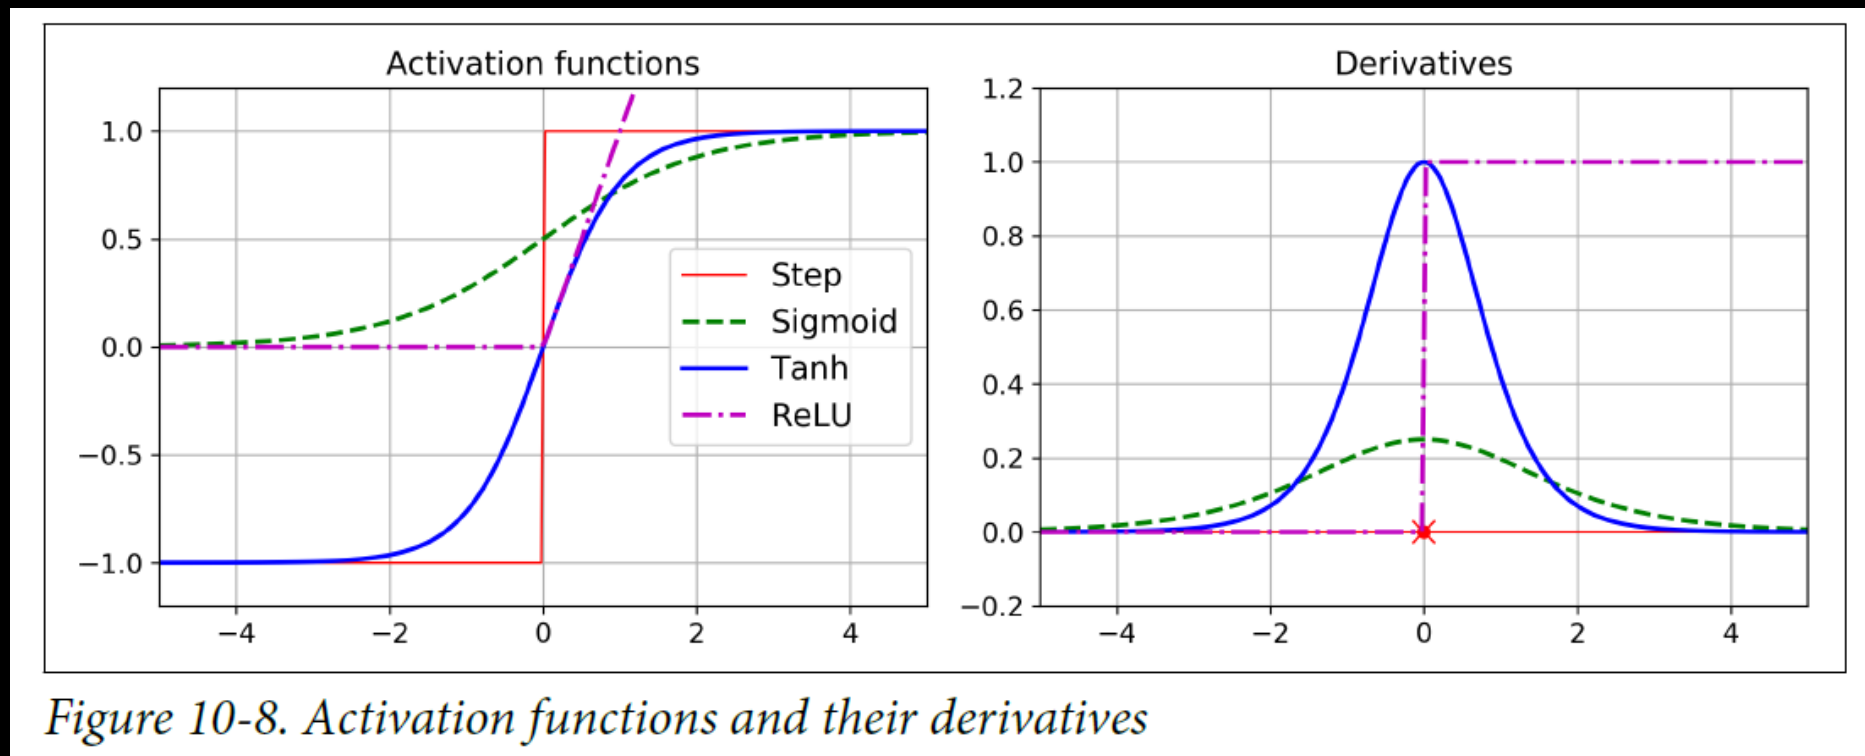


Backward pass



nervana

- Initialize connection weights randomly
- Use differentiable activation function
 - Sigmoid
 - Hyperbolic tangent (tanh)
 - Rectified Linear Unit (ReLU)



MLP for Regression

Table 10-1. Typical Regression MLP Architecture

Hyperparameter	Typical Value
# input neurons	One per input feature (e.g., $28 \times 28 = 784$ for MNIST)
# hidden layers	Depends on the problem. Typically 1 to 5.
# neurons per hidden layer	Depends on the problem. Typically 10 to 100.
# output neurons	1 per prediction dimension
Hidden activation	ReLU (or SELU, see Chapter 11)
Output activation	None or ReLU/Softplus (if positive outputs) or Logistic/Tanh (if bounded outputs)
Loss function	MSE or MAE/Huber (if outliers)

MLP for Classification




Figure 10-9. A modern MLP (including ReLU and softmax) for classification

Table 10-2. Typical Classification MLP Architecture

Hyperparameter	Binary classification	Multilabel binary classification	Multiclass classification
Input and hidden layers	Same as regression	Same as regression	Same as regression
# output neurons	1	1 per label	1 per class
Output layer activation	Logistic	Logistic	Softmax

Hyperparameter	Binary classification	Multilabel binary classification	Multiclass classification
Loss function	Cross-Entropy	Cross-Entropy	Cross-Entropy

Fine-Tuning NN Hyperparameters

- A lot of hyperparameters to tweak
 - Grid Search and Randomized Search
- Many companies offer services for hyperparameter optimization
- Active area of research
 - Evolutionary algorithms
- Some guidelines on:
 - Number of hidden layers
 - Number of neurons per hidden layer
 - Other hyperparameters

Number of Hidden Layers

- Start with a single hidden layer, gradually increase the number of layers. Stop just before overfitting
- For complex problems, higher number of hidden layers may be needed
 - Lower layers model low-level structures (line segments, shapes, orientations)
 - Intermediate layers model intermediate structures (squares, circles)
 - Higher layers model high-level structures (faces)
- Transfer learning
 - Start with pre-trained lower layers
 - Train higher layers with domain specific data

Number of Neurons per Hidden Layer

- Data determines the size of input and output layers
 - For MNIST, $28 \times 28 = 784$ input neurons and 10 output neurons
- Common practice: form a *pyramid* with fewer neurons at each layer
 - No longer standard practice
- Finding the perfect number of neurons is not well understood
- Stretch pants approach
 - Build a model with more layers and neurons than needed
 - Use early stopping or dropout to avoid overfitting

Other Hyperparameters

- Learning Rate
 - Start with a large value, then divide by 3 and repeat
- Batch Size
 - Not too large, not too small
 - Typically smaller than 32 but greater than 10
- Activation Function
 - Good choice to choose ReLU activation function as default activation function for all hidden layers.
 - Use early stopping